

NSA Solid-State Physicist Interview Prep

Quantum Computing Technical Program Manager / Solid-State Physicist

Fort George G. Meade Complex, Maryland

GG-13/1 to GG-15/10; Senior to Expert level

Purpose

This document is an interview preparation packet based on the provided NSA/LPS job posting screenshots and the candidate background discussed previously: Ph.D. experimental nuclear physics, radiation effects, multiphysics simulation, heat/energy transport, microelectronics, Geant4, HPC, and TS/SCI-polygraph cleared work. It is intended for interview practice only. Do not copy these answers into any application narrative that requires certification that no AI or consultant was used.

Prepared for: Chan Wook (Luke) Kim, Ph.D.
June 9, 2026

Contents

1	High-level read of the job	2
1.1	What the interviewer is likely screening for	2
1.2	Your best positioning	2
1.3	One-minute opening pitch	2
2	Core technical themes to be ready for	2
3	Technical interview questions and model answers	3
3.1	Solid-state quantum devices and qubit fundamentals	3
3.2	Materials, defects, radiation, and energy transport	5
3.3	Modeling, simulation, and scientific judgment	8
4	Behavioral and leadership questions with model answers	10
5	Hard questions you should expect	14
6	Questions you can ask the interviewers	15
7	Interview strategy	15
7.1	Do say	15
7.2	Do not say	16
8	Final compact answer bank	16

High-level read of the job

Role thesis

This is not merely a generic “solid-state physicist” position. The posting reads like a hybrid role: deep quantum-hardware physics plus technical program leadership for LPS research investments. The strongest theme is understanding and reducing the fundamental limitations of superconducting qubits or solid-state spin qubits, with relevance to fault-tolerant quantum computing.

What the interviewer is likely screening for

- Can you reason physically about qubit limitations, not just recite quantum buzzwords?
- Can you connect material defects, interfaces, electromagnetic environments, packaging, readout, and control electronics to qubit performance?
- Can you lead or shape a research portfolio across academia, industry, government, and national laboratories?
- Can you communicate technical risk, research gaps, maturity levels, and investment priorities to leadership?
- Are you honest about the boundary between your direct experience and adjacent expertise?

Your best positioning

Candidate positioning statement

I should present myself as a Ph.D. physicist who has spent the last several years modeling radiation effects, energy deposition, defect evolution, heat transport, and device-level degradation in microelectronics. My strongest fit is not that I have already spent a decade building superconducting qubits. My strongest fit is that I understand how materials, defects, transport, radiation, thermal environments, and multiphysics coupling limit solid-state devices. Those mechanisms are directly relevant to qubit lifetime, coherence, packaging, reliability, and mission-relevant deployment.

One-minute opening pitch

I am a Ph.D. experimental nuclear physicist and computational physicist specializing in radiation effects, particle and energy transport, defect formation, heat transport, and multiphysics modeling for microelectronics. In my current work, I build simulation workflows that connect incident particles and photons to energy deposition, material response, and device-level degradation. That background maps naturally onto the LPS mission of understanding fundamental limitations in solid-state quantum devices, because qubit performance is ultimately limited by materials, interfaces, defects, thermalization, packaging, electromagnetic control, and environmental disturbances. I would bring a rigorous physics-based modeling approach, experience in classified R&D environments, and the ability to connect microscopic mechanisms to system-level performance and research strategy.

Core technical themes to be ready for

Theme	What they may ask	How to connect your background
Superconducting qubits	Transmons, Josephson junctions, TLS loss, quasiparticles, readout, packaging	Your bridge: materials defects, energy deposition, phonon/thermal transport, interface loss, radiation-induced quasiparticles
Spin qubits	Quantum dots, gate-defined confinement, charge noise, spin readout, valley physics	Your bridge: solid-state device physics, electrostatics, defects, charge trapping, transport modeling
Materials limitations	Defects, interfaces, dielectrics, two-level systems, surface loss	Your bridge: defect evolution, damage modeling, material response, microelectronics degradation
Environmental effects	Radiation, cosmic rays, phonons, thermal spikes, magnetic/electric noise	Your bridge: radiation effects, Geant4, energy deposition, heat transport, shielding and mitigation
QEC and fault tolerance	Error thresholds, physical error rates, correlated errors, leakage	Your bridge: translate microscopic failure mechanisms into device/system risk metrics
Technical program management	Portfolio strategy, research gaps, milestones, evaluation metrics	Your bridge: module-based simulation, test correlation, cross-functional integration, customer-facing roadmap work

Technical interview questions and model answers

Solid-state quantum devices and qubit fundamentals

Technical Q1

Question. What is a transmon qubit, physically?

Answer. A transmon is a superconducting qubit built from a Josephson junction shunted by a relatively large capacitance. The Josephson junction provides a nonlinear inductance, which makes the oscillator energy levels anharmonic. That anharmonicity allows one to isolate the lowest two levels as a qubit while avoiding uncontrolled excitation of higher states. Compared with the original charge qubit, the transmon operates in the regime where the Josephson energy is much larger than the charging energy, which makes it less sensitive to charge noise. The tradeoff is reduced anharmonicity, so pulse design and leakage control become important.

Technical Q2

Question. What is a Josephson junction, and why is it essential in superconducting qubits?

Answer. A Josephson junction is usually two superconductors separated by a thin insulating barrier. Cooper pairs tunnel through the barrier coherently, giving a current-phase relation rather than an ordinary linear resistor response. In circuit terms, it behaves like a nonlinear, nondissipative inductive element. That nonlinearity is essential because a purely linear LC oscillator has equally spaced levels and cannot define an addressable two-level qubit. The junction creates the needed anharmonicity.

Technical Q3

Question. What limits coherence in superconducting qubits?

Answer. The main limitations include dielectric loss from interfaces and amorphous oxides, two-level system defects, quasiparticle poisoning, flux noise, charge noise, Purcell loss through the readout environment, radiation-induced events, packaging modes, crosstalk, and control/readout imperfections. I would classify the mechanisms into material loss, electromagnetic-environment loss, nonequilibrium quasiparticle/phonon effects, and control-system errors. My background is strongest in the material, radiation, and energy-transport side: connecting external perturbations to microscopic excitations and then to device-level degradation.

Technical Q4

Question. What is a two-level system defect and why does it matter?

Answer. A two-level system, or TLS, is a microscopic defect or configurational degree of freedom that can occupy two nearly degenerate states and couple to the qubit electric field or strain field. In superconducting circuits, TLS defects in dielectrics, oxides, and interfaces can absorb energy from the qubit, generate noise, and create avoided crossings or time-dependent fluctuations. They matter because superconducting qubits store energy in electromagnetic fields that strongly sample surfaces and interfaces. Even a small density of lossy defects can limit T1 and frequency stability.

Technical Q5

Question. How would you investigate whether TLS loss is material-driven or geometry-driven?

Answer. I would separate participation ratio from intrinsic loss. First, simulate the electromagnetic field distribution and compute how much electric-field energy participates in each material region: substrate, metal-air interface, substrate-air interface, metal-substrate interface, dielectrics, and packaging. Then fabricate or compare devices where geometry changes the participation ratios while materials stay fixed, and devices where materials/processes change while geometry is held fixed. If loss tracks participation ratio, geometry and field concentration are dominant. If loss changes strongly with process at fixed participation, the intrinsic material loss tangent or defect density is likely changing.

Technical Q6

Question. What is quasiparticle poisoning?

Answer. In a superconductor, the ground state consists of Cooper pairs. Energetic disturbances can break Cooper pairs and create quasiparticles. These nonequilibrium quasiparticles can tunnel, absorb energy, change parity, and produce decoherence or relaxation in superconducting qubits. Sources can include stray infrared radiation, cosmic rays, radioactive contamination, thermal radiation, microwave leakage, or high-energy particle interactions. My radiation-effects background is directly relevant here because the question becomes: what deposits energy, where does it go, how does it downconvert into phonons and quasiparticles, and how long do those excitations survive?

Technical Q7

Question. How could radiation affect superconducting qubit arrays?

Answer. Radiation can deposit energy in the substrate, metal, or nearby package. That en-

ergy can create high-energy phonons, break Cooper pairs, generate quasiparticles, produce charge offsets, or excite spurious modes. A key concern is correlated error: one high-energy event may disturb multiple qubits over a spatial region, which is more dangerous for quantum error correction than independent single-qubit errors. I would model the chain as incident particle or photon transport, energy deposition, phonon/quasiparticle generation, diffusion and recombination, coupling to resonators/qubits, and finally the induced error signature.

Technical Q8

Question. How do spin qubits differ physically from superconducting qubits?

Answer. A spin qubit stores information in the spin degree of freedom of an electron or hole, often confined in a semiconductor quantum dot. Superconducting qubits encode information in collective circuit degrees of freedom, such as charge/phase across a Josephson junction. Spin qubits can be compact and potentially CMOS-compatible, but they are sensitive to charge noise through gate-defined confinement, nuclear spin noise depending on material, valley physics in silicon, and disorder at interfaces. Superconducting qubits are larger but easier to couple and control at present. Both are ultimately limited by solid-state materials and noise.

Technical Q9

Question. What is a gate-defined quantum dot?

Answer. A gate-defined quantum dot is a small confinement potential created electrostatically in a semiconductor heterostructure or MOS structure. Metallic gates shape the local potential landscape so that one or a few carriers are trapped in a small region. The dot can behave like an artificial atom with discrete energy levels. For spin qubits, the relevant states can be electron spin states, controlled by microwave or electric-dipole spin resonance and measured through spin-to-charge conversion.

Technical Q10

Question. What physical mechanisms limit spin qubit performance?

Answer. Important mechanisms include charge noise, valley splitting variability, interface roughness, hyperfine coupling to nuclear spins, spin-orbit coupling effects, phonon-induced relaxation, disorder, gate leakage, crosstalk, and readout infidelity. In silicon spin qubits, isotopic purification reduces nuclear spin noise, but interface and charge-noise problems remain central. I would pay particular attention to how defects, traps, and gate-stack materials create fluctuating electric fields that perturb the dot potential and exchange coupling.

Materials, defects, radiation, and energy transport

Technical Q11

Question. How would you connect defect physics to qubit performance?

Answer. I would map the defect to the physical observable it perturbs. A charged defect can shift electrostatic potential, create charge noise, and perturb spin-qubit confinement. A dielectric TLS can absorb electromagnetic energy and limit superconducting-qubit T₁. A magnetic impurity can generate spin noise. A structural defect can scatter phonons and alter energy relaxation. The key is to avoid saying “defects are bad” generically. The useful model is: defect type, location, coupling

channel, spectral density, time scale, and effect on measurable qubit parameters such as T1, T2, frequency drift, readout error, or correlated-error rate.

Technical Q12

Question. What is the difference between T1 and T2?

Answer. T1 is the energy relaxation time: how quickly the excited-state population decays to the ground state. T2 is the coherence or dephasing time: how quickly the relative phase of a superposition is lost. T2 includes both energy relaxation and pure dephasing. In simple cases, $1/T_2 = 1/(2T_1) + 1/T_\phi$, where T_ϕ is the pure dephasing time. Different physics can dominate each: dielectric loss or quasiparticles often affect T1, while low-frequency charge or flux noise strongly affects dephasing.

Technical Q13

Question. How would you model radiation-induced quasiparticle generation?

Answer. I would start with a particle-transport calculation to obtain spatial and temporal energy deposition. Then I would model the downconversion of deposited energy into high-energy phonons and quasiparticles, including phonon propagation, pair-breaking, quasiparticle diffusion, trapping, recombination, and coupling to the superconducting resonators or qubits. For a practical first model, I would use a reduced source term from Geant4, a diffusion-recombination equation for quasiparticle density, and compare predicted relaxation events or burst rates against measured T1 drops or simultaneous error events.

Technical Q14

Question. What is the value of Geant4 for this problem?

Answer. Geant4 is useful for simulating how photons, electrons, protons, neutrons, muons, or other particles interact with detector or device geometries and where they deposit energy. For quantum hardware, Geant4 can answer where external radiation deposits energy, what secondary particles are produced, and how shielding or packaging changes the energy-deposition distribution. Geant4 alone does not solve qubit decoherence. Its output has to be coupled to materials, phonon, quasiparticle, electromagnetic, and device-response models.

Technical Q15

Question. How would you validate a simulation workflow for quantum-device radiation response?

Answer. I would validate in layers. First, validate particle transport and energy deposition against known stopping powers, attenuation data, or test structures. Second, validate thermal or quasiparticle transport using controlled injection experiments, temperature dependence, or resonator response. Third, validate device response using measured changes in T1, T2, frequency shifts, burst rates, or correlated-error signatures under known radiation environments. I would avoid claiming a full predictive model until each layer has been benchmarked.

Technical Q16

Question. What are correlated errors, and why are they dangerous for fault tolerance?

Answer. Quantum error correction assumes errors are sufficiently sparse and often approximately local or independent. A correlated error occurs when one physical event causes multiple qubits to fail together, spatially or temporally. Radiation-induced phonon or quasiparticle bursts are an example: one event can create a cloud of excitations affecting many devices. Correlated errors are dangerous because they can violate the assumptions of the error-correcting code and produce logical errors at rates not predicted by independent-error models.

Technical Q17

Question. How would you distinguish thermal effects from direct electromagnetic or charge effects in a qubit experiment?

Answer. I would vary the relevant time scales, spatial layout, and shielding. Thermal or phonon-mediated effects often have propagation and relaxation time signatures and may affect multiple devices with delay or spatial gradients. Direct electromagnetic pickup may track wiring, filtering, resonances, and control lines. Charge effects may appear as frequency jumps, offset-charge shifts, or telegraph noise. A good experiment uses synchronized event detection, multiple sensors, control resonators, temperature dependence, and changes to packaging/filtering to isolate the mechanism.

Technical Q18

Question. What role does packaging play in qubit performance?

Answer. Packaging defines the electromagnetic, thermal, mechanical, and radiation environment of the chip. It can introduce spurious modes, slotline modes, seam loss, thermalization bottlenecks, infrared leakage, magnetic impurities, stress, and unwanted coupling paths. Good packaging is not an afterthought; it is part of the quantum device. For a government lab concerned with future quantum systems, packaging is also where laboratory physics becomes deployable hardware.

Technical Q19

Question. What is your view of interface loss?

Answer. Interface loss is often dominant because electric fields in superconducting circuits concentrate near metal-air, substrate-air, and metal-substrate interfaces. Even if the lossy region is only a few nanometers thick, its field participation can be significant. The right way to reason about it is through participation ratio times intrinsic loss tangent. This is why fabrication, surface treatment, substrate choice, metal deposition, cleaning, and geometry all matter.

Technical Q20

Question. How would you prioritize research into qubit-material limitations?

Answer. I would prioritize based on impact, controllability, and testability. First, identify which loss or noise mechanisms dominate current devices. Second, estimate which mechanisms will become limiting as scale increases. Third, choose experiments that can discriminate mechanisms rather than merely improve one device. Fourth, connect each research line to measurable outcomes: T1, T2, frequency stability, gate fidelity, readout fidelity, correlated-error rate, yield, reproducibility, or system integration risk.

Modeling, simulation, and scientific judgment

Technical Q21

Question. How do you decide how much physics to include in a model?

Answer. I start from the decision the model must support. If the decision is shielding thickness, the model may initially need accurate energy deposition and attenuation. If the decision is qubit layout, the model may need spatial propagation of phonons or quasiparticles. If the decision is materials processing, defect chemistry and interface loss matter more. I prefer modular modeling: start with the simplest falsifiable model, validate it against data, then add physics only when it changes the prediction or decision.

Technical Q22

Question. What is a reduced-order model and why might it matter here?

Answer. A reduced-order model compresses a more detailed physical model into a faster surrogate while preserving the outputs needed for decisions. For quantum-device research portfolios, this matters because full multiphysics simulations may be too slow for parameter sweeps, design optimization, or uncertainty analysis. I would use detailed simulations to generate trusted training or calibration data, then construct reduced models for rapid screening of geometry, materials, shielding, or operating conditions.

Technical Q23

Question. How would you handle uncertainty in your simulations?

Answer. I would distinguish numerical uncertainty, model-form uncertainty, parameter uncertainty, and experimental uncertainty. For example, uncertainty may come from unknown defect densities, material loss tangents, boundary conditions, source spectra, or thermalization rates. I would propagate uncertainties to decision-relevant outputs and identify which inputs dominate the prediction. This tells us what experiments would most reduce uncertainty.

Technical Q24

Question. What is the difference between verification and validation?

Answer. Verification asks whether we solved the equations correctly. Validation asks whether the equations and assumptions represent reality for the intended use. In practice, verification includes mesh convergence, time-step convergence, conservation checks, and comparison with analytic limits. Validation requires comparison with experiments or trusted data. For this role, the distinction is important because quantum hardware is sensitive to small mechanisms; a code can be numerically correct but physically incomplete.

Technical Q25

Question. How would you evaluate whether a proposed quantum technology is mature?

Answer. I would look beyond a single impressive device metric. Maturity requires reproducibility, yield, scalability, control and readout integration, packaging, error characterization, environmental robustness, fabrication repeatability, credible path to QEC, and demonstrated improvement under realistic operating conditions. For government investment, I would ask what technical risk remains,

whether it is fundamental or engineering-limited, what milestone would falsify the approach, and what unique national capability the investment builds.

Technical Q26

Question. How would you define a good research milestone for this position?

Answer. A good milestone should be measurable, technically meaningful, and tied to a decision. For example: demonstrate a factor-of-two reduction in correlated-error burst rate under a controlled radiation environment using a specified mitigation; or show that a surface-treatment process reduces interface participation-normalized loss across multiple wafer runs. A weak milestone would be “study decoherence” because it does not define success or a decision point.

Technical Q27

Question. What would you do in your first 90 days?

Answer. First, I would map the current portfolio: platforms, performers, metrics, known bottlenecks, and decision timelines. Second, I would identify where my expertise in radiation, defects, transport, and multiphysics can contribute quickly. Third, I would meet internal and external stakeholders to understand what data exist and what cannot be discussed broadly. Fourth, I would propose one or two focused technical thrusts, such as radiation-induced correlated errors or materials/interface loss modeling, with clear validation experiments.

Technical Q28

Question. How would you compare superconducting qubits and spin qubits as investment areas?

Answer. Superconducting qubits are relatively mature in control, coupling, and processor demonstrations, but face challenges in coherence, wiring, packaging, cryogenic overhead, and correlated errors. Spin qubits offer compactness and possible semiconductor-manufacturing compatibility, but face challenges in variability, control complexity, readout, charge noise, and scaling uniformity. I would not frame it as one universally winning. The investment question is which bottlenecks are fundamental, which are engineering, and where government investment can remove a national capability gap.

Technical Q29

Question. How do you think about quantum error correction from a hardware perspective?

Answer. From a hardware perspective, QEC converts physical device imperfections into a statistical problem, but only under assumptions about error rates, locality, leakage, measurement fidelity, and correlation. Hardware development should therefore not optimize only average single-qubit gate fidelity. It must characterize error mechanisms in a way that is meaningful for codes: correlated errors, leakage, crosstalk, measurement errors, drift, burst events, and non-Markovian noise. My background is useful because radiation and materials failures can produce exactly the non-ideal error structures that are dangerous for QEC.

Technical Q30

Question. What is your answer if they ask, “You are not a qubit experimentalist. Why are you a fit?”

Answer. I would answer directly: I am not going to overstate my direct experience. My core expertise is in the physics that often limits advanced solid-state devices: radiation interactions, energy deposition, defect generation, transport, thermal response, and device degradation. Qubit platforms are also solid-state devices whose performance depends on materials, interfaces, noise, and environmental coupling. This role needs people who can connect microscopic physics to system-level quantum performance and research strategy. That is where I can contribute immediately while deepening platform-specific expertise.

Behavioral and leadership questions with model answers**Behavioral Q1**

Question. Tell us about yourself.

Answer. I am a Ph.D. experimental nuclear physicist who has moved into computational physics for radiation effects and microelectronics. My work has focused on connecting microscopic interactions, such as particle or photon energy deposition, to materials response, heat transport, and device-level consequences. I have worked in cleared defense environments, so I understand the importance of mission focus, technical rigor, and communication across teams. This position interests me because solid-state quantum devices are limited by the same type of physics I have spent years studying: materials, defects, transport, energy flow, and system-level vulnerability.

Behavioral Q2

Question. Why NSA/LPS?

Answer. LPS is attractive because it sits at the boundary between fundamental physics and national need. The role is not just to publish interesting quantum-device papers; it is to identify what technical barriers matter for future U.S. computing capability and direct research toward those barriers. That combination of deep physics, classified mission relevance, and collaboration with universities, national labs, and industry is exactly where I think my background can be useful.

Behavioral Q3

Question. Why this role instead of staying in radiation effects?

Answer. I do not see it as abandoning radiation effects. I see it as applying my strongest physics to a more advanced device platform. Quantum hardware is vulnerable to materials defects, radiation, quasiparticles, thermalization bottlenecks, electromagnetic environments, and packaging limitations. Radiation effects and transport modeling are part of the broader problem of understanding qubit limitations. This role would let me apply that expertise to a national-priority technology area.

Behavioral Q4

Question. Describe a time you had to learn a new technical area quickly.

Answer. In my current work, I had to connect radiation transport to device-level degradation,

which required understanding particle interactions, semiconductor response, heat transport, and numerical modeling. I approached it by decomposing the problem into modules: source and energy deposition, defect response, electrostatics, thermal response, carrier transport, and device metrics. That modular strategy let me learn quickly while still producing a coherent workflow. I would use the same method here for platform-specific qubit physics: identify mechanisms, equations, measurements, and decision points.

Behavioral Q5

Question. Tell us about a time you communicated complex physics to non-specialists.

Answer. A recurring part of my work has been translating simulation results into engineering or program decisions. For example, instead of presenting only particle-transport details, I explain how energy deposition changes material response, how that maps to device-level risk, and what design or shielding options follow. I try to separate the physics mechanism, the evidence, the uncertainty, and the recommended action. That style is important for this role because quantum-device limitations have to be communicated to both specialists and leadership.

Behavioral Q6

Question. How do you handle disagreement with another technical expert?

Answer. I try to move the disagreement from opinion to discriminating evidence. First, I restate the technical disagreement clearly: which assumption, mechanism, or prediction differs? Second, I ask what measurement or calculation would distinguish the two views. Third, I look for a low-cost test or benchmark. In physics, the goal is not to win the argument; it is to converge toward the mechanism that best explains the data.

Behavioral Q7

Question. Describe a time your model was wrong.

Answer. The right answer is to be transparent. In complex multiphysics work, early models are often incomplete. When predictions do not match data, I examine whether the issue is numerical error, boundary conditions, missing physics, or incorrect parameters. I do not treat a mismatch as failure; I treat it as information about the system. The important thing is to update the model rigorously and document what changed and why.

Behavioral Q8

Question. How do you prioritize multiple technical tasks?

Answer. I prioritize by mission impact, decision urgency, technical risk, and dependency structure. If a quick calculation can eliminate a bad design path, I do that early. If a deeper model is needed for a major decision, I define milestones and validation checks. I also communicate what is known, what is uncertain, and when a result will be decision-quality rather than exploratory.

Behavioral Q9

Question. How would you lead a research portfolio?

Answer. I would start by defining the strategic objective and the key technical bottlenecks.

Then I would map each project to a bottleneck, metric, and decision point. I would avoid funding disconnected demonstrations that do not reduce uncertainty. Good portfolio leadership requires balancing near-term deliverables with long-term capability building, and making sure each performer is producing evidence that changes what we know.

Behavioral Q10

Question. How do you work with academia, industry, and government partners?

Answer. I would recognize that each partner has different incentives. Academic groups often optimize novelty and publications. Industry optimizes product roadmaps and proprietary advantage. Government labs optimize mission need and long-term capability. The program manager's role is to align those incentives around measurable technical questions, protect sensitive information, and ensure that results are comparable across performers.

Behavioral Q11

Question. Give an example of improving efficiency or outcomes.

Answer. One example is building modular simulation workflows rather than one-off analyses. By separating energy deposition, defect response, transport, and device metrics, the workflow becomes reusable across geometries and scenarios. That reduces repeated manual work, makes assumptions easier to audit, and allows faster design comparisons. For this role, I would bring the same approach to evaluating qubit-limiting mechanisms and mitigation options.

Behavioral Q12

Question. How do you make decisions with incomplete data?

Answer. I separate what is known, what is uncertain, and what would change the decision. Then I look for robust decisions: options that remain good under multiple plausible assumptions. If the uncertainty is decision-critical, I propose a targeted measurement or simulation to reduce it. In advanced R&D, waiting for perfect data can be as damaging as acting blindly. The right approach is disciplined uncertainty management.

Behavioral Q13

Question. How do you respond to classified or sensitive constraints on publication and collaboration?

Answer. I respect that mission constraints come first. At the same time, strong science depends on careful documentation, internal peer review, and appropriate external engagement when allowed. I would be deliberate about separating open scientific principles from sensitive applications, following review processes, and ensuring that collaborations protect both national security and technical integrity.

Behavioral Q14

Question. What is your leadership style?

Answer. My leadership style is evidence-driven and technically engaged. I do not think leadership in a research environment means simply assigning tasks. It means understanding the physics

well enough to ask the right questions, identify weak assumptions, and help the team converge toward decision-quality evidence. I also value clear documentation and honest communication about uncertainty.

Behavioral Q15

Question. What weakness should you acknowledge?

Answer. A fair answer is that my direct hands-on experience with superconducting-qubit fabrication and dilution-refrigerator measurements is not as deep as someone coming directly from a qubit experimental lab. But I would immediately add that my adjacent strengths are highly relevant: solid-state device physics, radiation effects, energy transport, defect mechanisms, simulation, and classified mission execution. I would also emphasize that I learn new physics quickly and do not overclaim expertise.

Behavioral Q16

Question. How do you handle high-stakes technical responsibility?

Answer. I rely on traceability: assumptions, inputs, equations, validation status, and uncertainty must be documented. In high-stakes environments, a result is not useful if no one can audit it. I also distinguish exploratory analysis from decision-quality analysis. That distinction prevents premature confidence and helps leadership understand what can be acted on.

Behavioral Q17

Question. Why should we choose you over a conventional quantum-computing candidate?

Answer. A conventional quantum candidate may have more direct qubit-lab experience. My differentiator is the ability to connect external environments, radiation, materials defects, heat transport, and device degradation to quantum-hardware performance. As quantum systems scale and move toward mission-relevant environments, those issues become more important. I would complement, not duplicate, a team of qubit specialists.

Behavioral Q18

Question. Where do you see yourself contributing most in the first year?

Answer. I would contribute most in mechanism-driven modeling and evaluation of qubit limitations, especially those related to radiation, materials defects, phonon/quasiparticle transport, thermalization, and device-level risk. I could also help structure research programs with clear metrics and validation plans. My goal in the first year would be to become a trusted technical bridge between microscopic physics and program-level decisions.

Behavioral Q19

Question. What motivates you about public service?

Answer. The most credible answer is mission plus technical responsibility. I am motivated by the chance to apply rigorous physics to problems that matter for national capability. In industry or academia, technical work can be excellent but may be disconnected from urgent national needs. This role offers a way to use deep scientific training in service of long-term U.S. technological leadership.

and security.

Behavioral Q20

Question. Do you prefer individual contributor work or program leadership?

Answer. I prefer a role that combines both. I want to remain technically deep enough to judge the physics myself, but I also want to influence research direction and connect technical work to mission outcomes. This posting is attractive because it requires both: understanding qubit-device limitations at the mechanism level and shaping research investments across a broader community.

Hard questions you should expect

Hard Q1

Question. Your resume is radiation effects and nuclear physics. Why should we believe you can work on qubits?

Answer. Because the limiting physics of qubits is not isolated from radiation, materials, defects, and transport. Superconducting and spin qubits are solid-state devices whose performance depends on interfaces, losses, charge traps, phonons, quasiparticles, thermalization, and packaging. My expertise is in modeling how microscopic physical interactions propagate into device-level effects. I would be honest that I need to deepen platform-specific details, but the underlying physics skill set is directly relevant.

Hard Q2

Question. Have you personally fabricated or measured superconducting qubits?

Answer. No, not as my primary work. My direct experience is in radiation effects, particle and energy transport, defect/device modeling, and microelectronics. I would not claim hands-on dilution-refrigerator qubit measurement expertise. The value I bring is adjacent and complementary: understanding how radiation, defects, materials, and transport mechanisms can limit qubit performance and how to model and mitigate those effects.

Hard Q3

Question. What would you need to learn fastest to succeed here?

Answer. I would prioritize three areas: first, the specific device architectures LPS is most focused on; second, the dominant measured failure modes and metrics for those platforms; and third, the existing performer landscape and research portfolio. Technically, I would deepen my knowledge of superconducting circuit Hamiltonians, spin-qubit control/readout, cryogenic packaging, and QEC-relevant error characterization.

Hard Q4

Question. How would you avoid becoming a purely administrative technical program manager?

Answer. I would stay tied to mechanism-level technical questions and measurable outcomes. For every program, I would want to know what physical limitation is being tested, what evidence would

support or falsify the approach, and how the result changes the investment decision. That keeps program management grounded in physics rather than schedules alone.

Hard Q5

Question. What is one research thrust you would propose?

Answer. I would propose a thrust on radiation and materials-driven correlated errors in superconducting or spin-qubit systems. The goal would be to connect environmental radiation and material energy relaxation pathways to QEC-relevant error statistics. The program would combine controlled radiation experiments, device arrays with event sensors, Geant4-based energy deposition, phonon/quasiparticle or charge-trap transport models, and mitigation through shielding, substrate engineering, phonon traps, quasiparticle traps, or packaging design.

Questions you can ask the interviewers

1. Which platform is currently the stronger focus for this role: superconducting qubits, spin-based gate-defined quantum dots, or both?
2. Is the role expected to be primarily hands-on technical research, technical program direction, or a balance of the two?
3. What are the most important performance limitations LPS is trying to understand: coherence, gate fidelity, readout, correlated errors, packaging, fabrication variability, or scaling?
4. How does LPS define success for this position in the first year?
5. How much of the role involves coordinating external performers versus internal technical analysis?
6. Are radiation effects, quasiparticle bursts, cryogenic packaging, or materials defects active areas of interest in this portfolio?
7. What kinds of technical backgrounds are already strong on the team, and where is there a gap?
8. How are research investments evaluated: device metrics, QEC relevance, transition potential, national capability, or all of these?
9. How does LPS balance open scientific collaboration with classified mission needs?
10. What would make a candidate successful at the GG-15 or Expert level in this role?

Interview strategy

Do say

- “My direct qubit-lab experience is limited, but my physics background is highly relevant to qubit limitations.”
- “I connect microscopic mechanisms to device-level and system-level performance.”
- “Radiation, defects, phonons, quasiparticles, interfaces, and packaging become more important as quantum systems scale.”
- “I think in terms of falsifiable mechanisms, validation data, uncertainty, and decision-quality metrics.”

- “I can contribute as a technical bridge between physics mechanisms and research strategy.”

Do not say

- Do not claim to be an expert in superconducting-qubit fabrication if that is not true.
- Do not overfocus on nuclear physics unless you connect it to quantum-device limitations.
- Do not imply that radiation effects alone explain all qubit limitations.
- Do not sound like you want this only as a path to academia.
- Do not give generic answers about “quantum is the future” without mechanism-level detail.

Final compact answer bank

Thirty-second answer: strongest fit

My strongest fit is the physics of device limitation. I have deep experience with radiation interactions, energy deposition, defect response, heat transport, and microelectronics degradation. Superconducting and spin qubits are solid-state devices whose performance is limited by materials, interfaces, defects, noise, thermalization, and environmental coupling. I can help LPS connect those microscopic mechanisms to qubit performance and research investment decisions.

Thirty-second answer: gap

The main gap is that I am not coming directly from a superconducting-qubit or spin-qubit experimental lab. I would be transparent about that. But I bring adjacent depth in radiation effects, transport, materials response, device modeling, and cleared mission work. I think that makes me complementary to a team of platform specialists.

Thirty-second answer: why this matters

As quantum systems scale, rare environmental events, correlated errors, materials variability, packaging losses, and defect-driven noise become increasingly important. Those are exactly the kinds of problems that require solid-state physics, transport modeling, radiation understanding, and system-level thinking.